

Town Crier: An Authenticated Data Feed for Smart Contracts

Your N. Here
Your Institution

Second Name
Second Institution

Name
Name Institution

Abstract

Smart contracts are programs that execute autonomously on blockchains. Many of their envisioned uses require them to consume data from outside the blockchain. (e.g., financial instruments rely on equity prices.) Trustworthy data feeds that can interface with smart contracts will thus be critical to any smart-contract system.

We present an authenticated data feed system called Town Crier (TC). TC builds on the observation that many web sites, such as major news and finance sites, already serve as trusted data sources for non-blockchain uses. TC acts as a bridge between such servers and smart contract systems, using trusted hardware to authenticate and scrape data from HTTPS-enabled websites and to generate authenticated, timestamped datagrams for relying smart contracts. It also includes a range of advanced features, such as private datagrams, which decrypt and evaluate ciphertext requests within TC’s hardware.

We describe the TC architecture, its underlying trust model, and its applications, and report on an implementation that uses the newly released SGX software development kit and furnishes data for the smart-contract system Ethereum. We will soon be launching TC as an on-line public service.

1 Introduction

2 Background

2.1 SGX

Intel’s Software Guard Extensions (SGX) is a set of new instructions that confer hardware protections on user-level code. Its goal is to provide *isolated execution*. SGX enables a process to execute in a protected address space known as an *enclave*. It protects the confidentiality and integrity of a process in an enclave from other software on the same host, including the operating system, as well

as from certain forms of hardware attack, such as memory probes.

An enclave process cannot make system calls, nor can it execute code outside the enclave region. As a means of communicating with processes outside the enclave, however, it can read and write memory outside the enclave region, consistent with OS setting of page permissions. Thus isolated execution in SGX may be viewed in terms of an ideal model in which a process is guaranteed to execute correctly and with perfect confidentiality, but relies on a (potentially malicious) operating system for network and file-system access. This model is a simplification, as SGX is known to expose some of the internal state of an enclave to the operating system (e.g., making page faults visible to the exception handler []), creating potential side-channel vulnerabilities. Nonetheless, in this paper, we assume the ideal model described above for isolated execution in SGX.

Another feature of SGX is its support for *attestation*, which allows a remote system to verify the software in an enclave and communicate securely with it. When an enclave is created, the CPU produces a hash of its initial state known as a *measurement*. The software in the enclave may at a later time request a report, which includes a measurement and any supplementary data provided by the process, such as a public key. This report may be digitally signed (by a trusted process called a “quoting enclave”) using a hardware-protected key to produce a proof that the measured software is running in an SGX-protected enclave. This proof, known as a *quote*, may be verified by a remote system. The associated public key can then be used to establish a secure channel with the enclave or verify signed data it emits. We use the generic term *attestation* to refer to a quote, and denote it by *att*. We assume that a trustworthy measurement of the code for the enclave component of TC is available to any client that wishes to verify an attestation. SGX signs quotes using a *group signature* scheme called EPID []. This is significant in our design of Town Crier because

EPID is a proprietary signature scheme and is not, e.g., supported in Ethereum.

SGX additionally provides a trusted time source via the function `sgx_get_trusted_time`. On invoking `sgx_get_trusted_time`, an enclave obtains a measure of time relative to a reference point labeled with a nonce. This reference point remains stable for a given nonce, but SGX does not provide a source of absolute or wall-clock time, a limitation that we must engineer around in TC.

2.2 HTTPS

2.3 Smart contracts

Smart contracts are the expression of contractual agreements, including financial instruments, as executable code. In the context of cryptocurrencies, the term refers specifically to autonomously executing scripts that reside on a blockchain and can manipulate control currency. Bitcoin has a scripting language that can serve to implement a limited form of smart contract, but it is not Turing-complete and lacks support for loops.

Ethereum is the first decentralized blockchain with a Turing-complete scripting language and thus full support for smart contracts. Other Turing-complete smart contract systems exist, such as Counterparty [], which runs as a Bitcoin overlay, but is not fully decentralized. Ethereum has its own associated cryptocurrency called *ether*. (At the time of writing, 1 ether has a market value of a little more than \$2 U.S.) While TC can be adapted in principle to any smart contract system, we report on an implementation directed at Ethereum.

A smart contract in Ethereum is represented as what is called a *contract account*, endowed with code, a currency balance, and persistent memory in the form of a key/value store. Contract code executes in response to receipt of a *message* from another contract or a *transaction* from a non-contract (*externally owned*) account. Thus, contract execution is always ultimately initiated by a transaction. Informally, a contract only executes when “poked,” and poking progresses through a sequence of entry points until no further message passing occurs (or until there is a shortfall in gas, as explained below).

A smart contract accepts messages as inputs to any of a number of designated functions. These entry points are determined by the contract creator and represent the API of the contract. Once created, a contract executes autonomously; it persists indefinitely, with even its creator unable to modify its code. (There’s one exception: a special opcode *suicide* will wipe code from a contract account.) As a simple abstraction, then, a smart contract may be viewed as an *autonomous agent* on the blockchain.

To prevent denial-of-service attacks or inadvertent in-

finite looping within contracts and in general to control resource expenditure by the network, Ethereum implements uses a resource called *gas* to power contracts. Op-codes in smart contracts have globally specified, fixed gas costs, as do the use of a contract’s persistent storage and the data in transactions and messages. Transactions and messages include a parameter (STARTGAS) specifying a bound on amount of gas expended by the computations they initiate. Gas is carried along the execution path induced when a transaction or message is passed to a contract, and depleted as instructions are executed or reads or writes are made to persistent storage. Should a function fail to complete due to a shortfall in gas, it is aborted and any state changes induced by the partial computation are rolled back to their pre-call state; previous computations along the call path, however, are retained.

Along with the STARTGAS parameter, a GASPRICE parameter is included that specifies the maximum amount in ether that the transaction is willing to pay per unit of gas. The transaction thus succeeds only if the initiating account has a balance of $\text{STARTGAS} \times \text{GASPRICE}$ ether and GASPRICE is high enough to be accepted by the system (miner).

The management of gas, as we show in our design of Town Crier can be delicate. Without careful construction, for example, the smart contracts representing TC’s interface on the Ethereum blockchain can be caused by an attacker to exhaust the ether used to power the delivery of datagrams.

2.4 Applications of ADFs for smart contracts

2.5 Basic terminology

We refer to a *smart contract* making use of the Town Crier service as a *relying contract*. In contexts where a relying contract has issued to TC a service request for a datagram, we call it a *requester*. We denote a requester by C_U . A party (or would-be party) to a relying contract, a person, organization, or server, is a *client*. Relying contracts—and requesters, by extension—are blockchain entities, while a client is an off-chain entity. A *data source*, or *source* for short, is an online server (running HTTPS) that provides data which TC uses to compose datagrams.

3 TC Architecture and Security Model

The Town Crier system includes three main components: The TContract (\mathcal{C}_{TC}), the Engine (\mathcal{E}), and the Relay (\mathcal{R}). The Engine and Relay reside on the TC server, while

the TContract resides on the blockchain. An architectural schematic showing the roles of these components is given in Figure 1.




Figure 1: **Basic Town Crier architecture.**

The TContract C_{TC} . The TContract (denoted by C_{TC}) is a smart contract that acts as the blockchain front end of the TC service. It is designed to present a simple API to a relying contract C_U for its requests from TC. Very simply, C_{TC} accepts datagram requests from a requester C_U and returns corresponding datagrams from TC. Additionally, C_{TC} manages TC monetary resources, which in Ethereum take the form of ether (money) and gas (“fuel” for contracts).

The Engine \mathcal{E} . The Engine ingests and fulfills datagram requests from the blockchain. To obtain the data for inclusion in datagrams, it queries external data sources, specifically HTTPS-enabled internet services. It returns a datagram to a requesting contract C_U as a digitally signed blockchain message. The Engine runs in an SGX enclave, and is thus secured against an adversarial OS as well as other process on the host.

The Relay \mathcal{R} . As an enclave process, the Engine lacks direct network access. Thus the Relay handles bidirectional network traffic on behalf of the Engine. Specifically, the Relay provides network connectivity from the Engine to three different types of entities:

1. *The Blockchain (the Ethereum system):* The Relay scrapes the blockchain in order to monitor the state of the TContract C_{TC} . In this way, it performs implicit message passing from C_{TC} to the Engine, as neither component itself has network connectivity. Additionally, the Relay places messages emitted from the Engine (datagrams) on the blockchain.
2. *Clients:* The Relay runs a web server to handle off-chain service requests from clients, specifically, requests for attestations from the Engine. As we soon explain, an attestation provides a unique public key for the Engine instance to the client and proves that

the Engine is executing correct code in an enclave and that its clock is correct in terms of absolute (wall-clock time). A client that successfully verifies an attestation can then safely create a relying contract C_U that uses the TC.

3. *Data sources:* The Relay relays traffic to and from data sources (HTTPS-enabled servers) queried by Engine.

The Relay is an ordinary user-space application. It does not benefit from integrity protection by trusted hardware and thus, unlike the Engine, can be subverted by an adversarial OS on the TC server, causing network delays or failures. A key design aim of TC, however, is that Relay should be unable to cause incorrect datagrams to be produced or users to lose fees paid to TC for datagrams (although they may lose gas used to fuel their requests). In general, the Relay *can only mount denial-of-service attacks against TC*.

Security model Here we give a brief overview of our security model for TC, providing more details in our security analysis in section ?? . We assume the following:

- *Engine security:* We make three assumptions about Engine: (1) Engine behaves honestly, i.e., correctly executes the TC protocol; (2) The private key sk_{TC} is known only the Engine; and (3) The Engine has an accurate (internal) real-time clock. (Specifically, the clock is accurate to within XXX, as we show experimentally.) We explain in the next section how we achieve these properties through use of SGX.
- *Blockchain communication.* Transaction and message sources are authenticable, i.e., a transaction m sent from an account \mathcal{P}_X (or message m from contract C_X) is identified by the receiving account as originating from X . Transactions and messages are integrity protected (as they are digitally signed by the sender), but not confidential.
- *Network communication.* The Relay (and other untrusted components of the TC server) can tamper with or delay communications to and from the Engine. (As we explain in our SGX security model, the Relay cannot otherwise observe or alter the behavior of the Engine.) Thus the Relay is subsumed by an adversary that controls the network.

4 TC Protocol

We now describe the operation of TC at the protocol level. First, we give a brief protocol overview. Then we describe and provide some formalism for specifying the

data flows in the system. Finally, in this section, we provide a component-level view of the protocol by specifying the functionalities embodied in the TContract, Relay, and Engine. We present these as ideal functionalities, inspired by the universal-composability (UC) framework, in order to abstract away implementation details and as a springboard for formal proofs of security. We omit details in this section on how payment is incorporated into TC; this rather complicated aspect of the system design is deferred to Section ??.

4.1 Datagram lifecycle

The lifecycle of a datagram may be briefly summarized in the following steps:

- **Initiate request.** C_U sends a datagram request to C_{TC} on the blockchain.
- **Monitor and relay.** The Relay monitors C_{TC} and relays incoming datagram requests to the Engine.
- **Securely fetch feed.** Based on params, the Engine contacts a data source via HTTPS and obtains the requested datagram. It forwards the datagram via the Relay to C_{TC} .
- **Return datagram.** C_{TC} checks the correctness of params and returns the datagram to C_U .

We now make this data flow more precise.

4.2 Data flows

A datagram request by C_U takes the form of a message $m_1 = (id, callback, params)$ to C_{TC} on the blockchain. Here, id is a unique request identifier (which we explain later how to compute in practice); $callback$ specifies the entry point in C_U to which the datagram is to be returned. (In principle, $callback$ could specify an entry point in a different contract, but TC does not yet adopt this generalization. $params$ specifies the requested datagram, e.g., $params := (url, spec, T)$, where url is the target data source and $spec$ specifies content of a the datagram to be retrieved (e.g., a stock ticker at a particular time), while T specifies the delivery time for the datagram.

C_{TC} forwards $m_2 = (id, params)$ to the Engine. It receives in return a return message $m_3 = (id, params, data)$ from the TC service, where $data$ is the datagram, i.e., contains the data (e.g., the desired stock ticker price). C_{TC} checks the consistency of params on the incoming and outgoing messages, and if they match forwards data to the entry point $callback$ in C_U in message m_4 .

Figure 2 shows the data flows involved in processing a datagram request. For simplicity, the figure omits the Relay, which is only responsible for data passing.

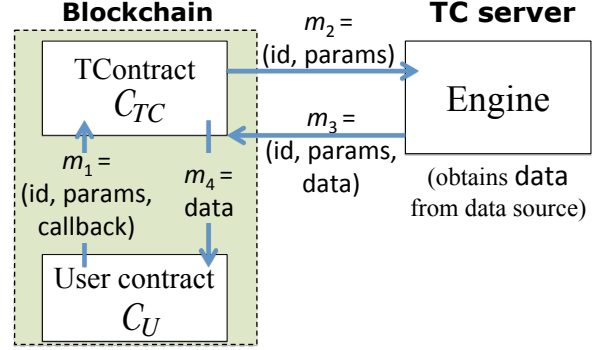


Figure 2: Data flows in datagram processing.

Digital signatures are needed to authenticated messages, such as ans, entering the blockchain from an external source. We let (sk_{TC}, pk_{TC}) denote the private / public keypair associated with the Engine for such message authentication. For simplicity, we assume for the time that the Engine can send signed messages directly to C_{TC} . Later we explain how Ethereum requires a slightly different approach in which TC sends messages via an Ethereum wallet.

4.3 Use of SGX.

Our protocols in TC rely on the ability of the SGX attestation mechanism to bind an enclave process instance to a public / private key pair and for the process to high-accuracy timestamps, in absolute or wall-clock time. To simplify our presentation, we defer most details of our formal model of SGX capabilities and specification of protocols for attestation generation to the upper appendix. Instead, we first explain and present a simple abstraction of these capabilities suitable for our TC protocol descriptions.

Let $prog_{encl}$ represent the code for Engine, which we presume is trusted by all system participants. To avoid having to verify an EPID group signature on the blockchain, we have clients obtain SGX attestations from the Relay and verify them off-chain. As noted above, it is possible to bind an enclave process instance running $prog_{encl}$ to a key pair (pk_{sgx}, sk_{sgx}) by including the public key pair in an attestation. We take this approach in TC. Specifically, Since Ethereum itself already verifies signatures on transactions from externally owned accounts (i.e., users interact with the Ethereum blockchain through an authenticated channel), TC uses a trick to *piggyback verification of enclave signatures on top of Ethereum's already existing transaction signature verification mechanism*. Very simply, the Engine associates pk_{sgx} with an externally owned Ethereum account P_{TC} . This approach saves not only software engineering effort,

but more importantly, the gas that would be required for EPID signature verification.

To make this idea work fully, the public key pk_{sgx} must be hardcoded into C_{TC} . A client creating or relying on a contract that uses C_{TC} is responsible for making sure that this hardcoded pk_{sgx} has an appropriate SGX attestation before interacting with the C_{TC} blockchain contract. Let $Verify$ denote a verification algorithm for EPID signatures. Figure 4.3 gives the protocol for a client to check that C_{TC} is backed by a valid Engine instance. This protocol does not include a mechanism for *revocation* of a compromised SGX instance, an issue we discuss later in the paper.

In summary, then, we may assume in our protocol specifications that *all relying clients have verified an attestation for Engine and thus that pk_{sgx} represents the public key of the associated instance.*

Clock. Additionally, as noted above, trusted clock provides only relative time with respect to a reference point, not absolute time. Thus, when initialized, the Engine is provided with the current wall-clock time by a trusted source, e.g., the Relay (under a trust-on-first-use model). The SGX attestation generated by the Engine includes the current wall-clock time, which clients may verify in real time. Thus, a client can determine the absolute clock time of Engine to within a high degree of accuracy, bounded by the round-trip time of its attestation request plus the attestation verification time—on the order of hundreds of milliseconds in a wide-area network [1]. [Ari: Let’s give the attestation API a function name.] We present a formal model and protocol specification in the paper appendix.

Notation. We model execution in SGX in terms of a functionality \mathcal{F}_{sgx} operating in a stateful manner on $prog_{encl}$. This functionality may be invoked through transmission of one of two messages to $prog_{encl}$: “initialize”, which creates the enclave with $prog_{encl}$ as its initial state and triggers measurement quotes and “resume”, which initiates an execution of $prog_{encl}$ on a fresh input. (We assume that $prog_{encl}$ exits only when it completes processing of a given input.) We let $\mathcal{F}_{sgx}[prog_{encl}, \mathcal{R}]$ denote invocation of \mathcal{F}_{sgx} on $prog_{encl}$ by \mathcal{R} . We let $clock()$ denote measurement of the SGX clock from within the enclave; $clock()$ returns the current wall-clock time (in a canonical format such as seconds and fraction sections since the Unix epoch January 1, 1970 00:00 UTC). [Ari: What format do we actually use?]

4.4 A payment-free basic protocol

For simplicity, we first specify functionalities in a payment-free version of our basic protocol, i.e., one that

User: offline attestation of SGX enclave

Inputs: $pk_{\mathcal{M}}, pk_{sgx}, prog_{encl}, \sigma_{att}$

Checks:

Assert $prog_{encl}$ is the expected enclave code
 Assert $\Sigma_{sgx}.Verify(pk_{\mathcal{M}}, \sigma_{att}, (prog_{encl}, pk_{sgx}))$
 Assert C_{TC} is correct and parametrized w/ pk_{sgx}
// now okay to rely on C_{TC}

Figure 3: A client checks an SGX attestation of the enclave’s code and its public key pk_{sgx} and checks that pk_{sgx} is hardcoded into TC blockchain contract C_{TC} before making use of C_{TC} .

does not include gas or fees—specifically the message component \$fee. Later, in our implementation discussion, we explain how we handle these two forms of payment, and we prove payment-related properties in the paper appendix.

Town Crier blockchain contract C_{TC}

Request: On recv (id, params, callback) from some contract C_U :

Record (id, params, callback)

Deliver: On recv (id, params, data) from pk_{sgx} :

Let (id, params, callback) be the most recently recorded tuple for id

Assert params = params’

Call callback(data)

Figure 4: A simple, fee-free version of the Town Crier contract C_{TC} . Note that communication with C_{TC} is through an authenticated channel implemented through digital signatures (which are not explicitly expressed in our notation).

The TContract

The Engine

The Relay

5 Implementation

5.1 TC Blockchain resources

These include the TC contract and the addresses from which it sends messages and manages its wallet

There are two parts to the Town Crier blockchain resources:

Town Crier's enclave program
Initialize: On receive “initialize”: $(pk_{sgx}, sk_{sgx}) := \Sigma.KeyGen(1^\lambda)$ Record (pk_{sgx}, sk_{sgx}) Output pk_{sgx} // incl. in measurement & attestation Resume: On receive (id, params) Parse params := (url, pk_{url} , T): $T_{start} := clock()$ Establish secure channel w/ url w/ public key pk_{url} Download the webpage at url $T_{end} := clock()$ Assert $round(T_{start}) = round(T_{end}) = T$ Parse webpage and extract data $\sigma := \Sigma.Sign(sk_{sgx}, (id, params, data))$ Output $((id, params, data), \sigma)$

Figure 5: SGX enclave's code.

Town Crier's untrusted relay \mathcal{R}
Initialize: Send “initialize” to $\mathcal{F}_{sgx}[prog_{encl}, \mathcal{R}]$ On receive (pk_{sgx}, σ_{att}) from $\mathcal{F}_{sgx}[prog_{encl}, \mathcal{R}]$: Publish (pk_{sgx}, σ_{att}) Loop forever: When \mathcal{C}_{TC} receives new request (id, -, params): Parse params := (url, pk_{url} , T) Fork: Wait till time T Send (“resume”, params) to $\mathcal{F}_{sgx}[prog_{encl}, \mathcal{R}]$ On recv $((id, params, data), \sigma)$ from $\mathcal{F}_{sgx}[prog_{encl}, \mathcal{R}]$: AuthSend (id, params, data) to \mathcal{C}_{TC}

Figure 6: Town Crier untrusted relay. For simplicity, here we assume that there is only a single enclave program. When multiple data feed sources are supported, we need multiple enclaves that instantiate different parsers for different sites. In this case, the Town Crier relay also initialize all enclave instances and route the request to the correct enclave instance.

- \mathcal{P}_{SGX} – An Ethereum wallet whose private key is generated and controlled by the SGX enclave.
- \mathcal{C}_{TC} – An Ethereum contract with two entry points, described in Table 12.

The Town Crier server is responsible for identifying calls to \mathcal{C}_{TC} 's Request request method by scraping the contract activity on the blockchain. It then replies to those requests through Deliver once it has acquired the necessary data and packaged it into a response.

The Town Crier contract \mathcal{C}_{TC}
Request: Upon receiving (type, callback, \$fee) from a user \mathcal{P} : If $(\$fee < F_{min} \text{ or } \$fee > F_{max})$ Return \$fee to \mathcal{P} Otherwise do nothing. Deliver: Upon receiving (callback, data, \$fee) from a user \mathcal{P} : If $\mathcal{P} \neq \mathcal{P}_{SGX}$ Return with no effect. Call callback(data) providing \$fee – F_{min} ether as the maximum gas. Send \$fee ether to \mathcal{P}_{SGX} .

Table 1: Definition of the Town Crier contract \mathcal{C}_{TC} . Here the minimum fee, F_{min} , is the amount of ether necessary to cover the gas costs of Deliver not including the execution of callback and the maximum fee, F_{max} , is the maximum amount of ether Town Crier can send as gas to a single execution of Deliver.

5.2 Client API

5.3 TC server

5.3.1 Trusted executable

5.3.2 Untrusted executable

6 Security Analysis

6.1 Gas neutrality

Here we assume an adversary which is active on the blockchain, the network, and within the untrusted executable running on the Town Crier server. However, we assume that the adversary will not execute an arbitrary denial of service attack, but will rather delay messages indefinitely and deliver bogus data whenever such data will be accepted as valid. Because operations on the blockchain are verifiable and the SGX enclave can attest to what it is running, we assume those are honest.

In this model we show that, for every request which provides a sufficient fee, a valid authenticated datagram will be delivered to the requested callback location in finite time. If the request includes an insufficient fee (but is otherwise valid), the datagram will not be delivered, but the (too-small) fee will still be collected.

Lemma 1. *If seeded with at least F_{max} ether, the \mathcal{P}_{SGX} wallet will have at least as much money after each transaction as it had before that transaction.*

Proof. [Ethan: This is actually a proof sketch, I just put it in a proof tag.]

Because all blockchain transactions from \mathcal{P}_{SGX} must be initiated by the SGX enclave and the SGX only calls $\mathcal{C}_{TC}.\text{Deliver}$, we need only reason about what happens inside that function. Because transactions including \mathcal{C}_{TC} are transmitted securely into the SGX enclave, it will only see valid requests (ones for which $F_{\min} \leq \$\text{fee} \leq F_{\max}$) and the arguments it sees for those requests will be correct. Moreover, it saves the transaction ID of each request it fulfills and never fulfills a request with the same transaction ID twice. This means that whenever deliver is called, it will be called in connection with a valid request that has not already been delivered. Thus it suffices to show that:

1. The first time a valid request is delivered, \mathcal{C}_{TC} will contain at least $\$ \text{fee}$ ether.
2. $\$ \text{fee}$ is never lower than the amount \mathcal{P}_{SGX} must spend in gas.
3. The execution of Deliver will never run out of gas (and thus always succeed).

To prove claim 1 we first note that ether can only be removed from \mathcal{C}_{TC} as part of a call to Deliver from \mathcal{P}_{SGX} . Because \mathcal{P}_{SGX} is honest, it will only make this call in connection with a valid request, and the specified value of $\$ \text{fee}$ will always be the fee submitted with that request. Because \mathcal{C}_{TC} pays out the specified $\$ \text{fee}$ on the call to Deliver, the ether is always exactly the ether stored from a previous, valid, undelivered request, and thus will always be present in \mathcal{C}_{TC} .

To prove claim 2 first note that a request is only considered valid if $\$ \text{fee} \geq F_{\min}$. F_{\min} is defined so that it is high enough to cover gas costs for all of Deliver except the execution of the provided callback. However, callback is only given $\$ \text{fee} - F_{\min}$ ether worth of gas to execute. Therefore it is impossible for the entire call of Deliver to spend more than $F_{\min} + (\$ \text{fee} - F_{\min}) = \$ \text{fee}$ ether on gas.

To prove claim 3 we note that $\$ \text{fee} \leq F_{\max}$ and, by construction, \mathcal{P}_{SGX} will always provide at least F_{\max} in gas for the execution of Deliver. Therefore we have that \mathcal{P}_{SGX} will always provide at least $\$ \text{fee}$ in gas to execute Deliver. By the argument above, Deliver can never use more than $\$ \text{fee}$ in gas, so therefore an SGX-initialized call to Deliver will never run out of gas. \square

Lemma 2. *Any data given as an argument to callback in \mathcal{C}_{TC} 's Deliver method is verifiably authentic.*

6.2 Other security concerns

We treat side-channel attacks as outside the scope of our initial TC architecture. Such attacks would be of particular concern should the Relay be compromised. Intel

explicitly disclaims protections against side-channel attacks in SGX. The ability for the OS to monitor page faults incurred by a process running in an enclave is an example shown to be potentially serious in practice []. Additionally, the Relay or any network adversary can potentially perform traffic analysis to determine what content the Engine is retrieving from a remote server [], a potential threat to the confidentiality of private datagrams.

7 Applications

Discuss flight insurance as an example: We'd like to conceal the flight number and date. We might also want to conceal payment, so TC might ingest encrypted addresses and mix them internally.

Micro-loans too? Linkage to Facebook / Keybase.io

8 Advanced Features

8.1 Custom and private datagrams

8.2 Full-blockchain scraping

As a means of reducing communication costs...

8.3 Use of new opcodes

New opcodes would enable processing of fresh attestation

8.4 Protecting against freeloading

8.5 Principled data extraction

Ultimately, we servers might themselves to act as ADFs. Possible migration path: (1) Town Crier; (2) XML labels on data; (3) Integration of Town Crier features into source directly

8.6 Off-chain communication

Mention use of Lamport signatures, etc.

8.7 Revocation

8.8 IoT support

9 Conclusion

A Introduction

B Identifying Client Contracts

The Authenticated Data Feed (ADF) needs some method to identify and serve prospective clients. There are two on-chain methods: registration and client flags. The registration method requires an additional transaction and will therefore incur an additional fee. The flags method must scan all contracts at every block update to support dynamic data requests. It may therefore be useful to support both methods, using the blockchain crawler for one-time requests and registration for more complicated contracts, depending on the resource costs of the crawler.

B.1 Registration / Explicit Requests for Data

This requires the client to initiate communication to the ADF with a *Initial Client Request* message, which consists of a (potentially zero-value) transaction to the ADF address with a message specifying what signed data it wants.

B.2 Client Flags / ADF Blockchain Crawler

This method requires only that the client contract has in its key/value store a flag indicating its request for service from the ADF and the specifics of the data it wants. The ADF can then crawl the blockchain looking for contracts with this specific flag set and read the data requested.

B.3 Off Chain Communication

The ADF may be notified of prospective clients by any manner of off-chain communication however there will be no public record of this request. The ADF may then selectively deny service. Malicious clients may also have an easier time flooding the ADF with requests, compared to the other two methods which have an Ether cost (contract creation).

B.4 Migration Path

Ultimately, we expect sources themselves to act as ADFs. The migration path is : (1) Town Crier; (2) XML labels on data; (3) Integration of Town Crier features into source directly

C Payment Methods

C.1 Separate Fair Exchange Contracts

The ADF or the client contract creates or uses an existing *fair exchange contract* for each data request. This contract requires input from the client in the form of some amount of Ether (this can happen during the creation of the contract), and the signed data from the ADF. Once both inputs have been received, the contract sends the Ether to the ADF and the data to the client. Note the client should send the currency first as once the ADF publishes the signed data it will be publicly visible (unless we use zero-knowledge proofs). If the ADF fails to deliver the data within a time period then the client's money is refunded.

The main benefit of this method is that the fair exchange contract should be easier to verify for both parties. Depending on the specifics of the attestation it needs to check, it may be possible for the contracts to be completely templated. Also note that creating a new contract will incur a transaction fee. These data fair exchange contracts can potentially batch data requests and verification and can be re-used for subsequent requests if desired. [Fan: the client and the ADF should agree on the content of this contract. In my mind this template could be provided by an ADF service provider along with other information on their website. For example, clients can browse ADF's website for pricing info, the public address, an template for generating request and ensuring fair exchange and the public key. Then clients can embed these information in there own contract.]

[[Kyle: This is my thought as well. A client may extract the address, pricing info, and a code template for fair exchange from the ADF website. The issue I was trying to get at was that verifying the bytecode of one function in a contract may be more difficult than simply verifying an entire contract's bytecode due to having to isolate the function. I don't know whether isolating a function is hard or if the bytecode may change depending on the client contract (offsets, compilers, etc), but at the very least it involves the extra step of checking the function entry point in addition to simply matching the hash of the bytecode]]

C.2 Payment Built-in to the Client Contract

The client contract includes a function that accepts signed data and pays out the specified amount of coins to the sender if the data is correct. This requires the ADF to verify the correctness of the function and that the client has sufficient funds. It would be potentially useful for this function to be templated for ease of use and verifi-

cation, and at the very least it needs to conform to some specification agreed upon by the client and ADF.

C.3 Payment built-in to the ADF

The client may send payment directly to an ADF contract as part of its initial request. The ADF contract code must maintain a list of clients and their corresponding requests. It must also verify the attestation on the data before forwarding it to the client since it has already been paid. The client must verify that the ADF contract code is correct.

[Fan: Do all of 9 combinations of (2.1, 2.2, 2.3) and (3.1, 3.2, 3.3) make sense? Or we should point out the best combinations?] [[Kyle: All combinations make sense to me with the exception of client flags and payment built-in to the ADF (2.2, 3.3)]]

D Privacy

Everything posted to the blockchain is publicly visible, including all data requests directed towards the ADF. It is possible to encrypt data requests, as the ADF may decrypt and process them off-chain. It is not always possible to encrypt the responses, as the contract requesting the data will often need to process the data.

D.1 Privacy From Peers

In some cases clients may want to keep their requests for data private. For example a client requesting information about the weather may need to provide his zip code for accurate local results. However he may not wish for his zip code to be publicly visible. To solve this, he may encrypt his zip code (or his entire request) with the ADF's public key before posting it to the blockchain. Note that all encryption and decryption operations can be done off-chain to avoid incurring excessive gas fees.

D.2 Partial Privacy From ADFs

Clients may also wish to keep some information private from the ADF. Consider again a client requesting information about the weather who wishes to keep his zip code private not only from the public, but from the ADF as well. This can be accomplished by utilizing two ADFs A_1 and A_2 . The client creates $n - 1$ randomized decoy requests for data along with his actual request. He orders these requests such that his actual request appears at a randomly chosen index i . He then encrypts all requests under the public key of A_1 and submits all n encryptions in order to A_1 . He encrypts i under the public key of A_2

and submits the encryption to A_2 . A_1 decrypts and computes the results for all n queries and submits the results to A_2 . A_2 returns to the client the i th result. So long as the two ADFs do not collude (in the context of the weather example, this means neither ADF can know both the zip codes and the index i), A_1 can only guess the clients true data with probability $\frac{1}{n}$. [[Kyle: There is definitely some leakage here if the response returned by A_2 is in plaintext]]

[[Kyle: The issue is that the information has to appear in plaintext somewhere on the blockchain in order for a contract to use it, otherwise A_2 could simply encrypt the response under the public key of the client. We could obfuscate the result with a mask, but the client contract will still need to unmask the data to use it, and when it does anyone can read it. Then in our example A_1 knows what the weather was at the clients zipcode and may cross reference with the zip-codes he was given]] [Ari: As regards leakage, that's right. If the output of the contract depends on the zip code, information will leak] [Ari: Remember the compression trick as well. Let G be an additive group over 5-decimal-digit numbers. Let $PRF_k : \{0, 1\}^* \rightarrow G$ be a PRF. To send a correct zip code z in list of n zip codes, of which $n - 1$ are "decoys", the client selects a key k and index $i \in [1, n]$ at random. The client sends $(k, p = PRF_k[i] - z)$. The ADF decodes the list as $\{PRF_k[1] - p, \dots, PRF_k[n] - p\}$.]

D.3 Fully Private Requests

If a client wishes to keep the entirety of the request private, he may do so by leveraging the ADF's trusted hardware. The initial request is encrypted under a public key stored in the trusted hardware, which processes the request. A malicious ADF operator may only determine the sources which are contacted as a result of the request, and if necessary these can be masked with decoy queries as well.

D.4 Decoy Requests

Decoy requests for partially private protocols should be in the same class as the real request. In the weather example, all zipcodes provided as decoy requests should be experiencing the same weather as in the actual request. As the outcome of the contract is public, the ADF will be able to determine the value of the weather at the client's zipcode and can discount all decoys that do not match.

If the class of decoy requests can be pseudorandomly generated, we can save data transmission fees by

sending only a random function, a seed, and a mask to ensure that the real request is computed.

E Example: Travel Insurance

An insurance company may provide insurance for canceled flights by maintaining an Ethereum contract on the blockchain. The source code of this contract should be public to allow potential customers to verify its correctness. A customer may then purchase insurance by paying a predetermined amount of ether into the contract. In addition, this transaction should include identifying information (flight number) for the flight for which the customer wishes to be insured. However, the customer may wish to hide the flight number in order to avoid publicly revealing which flight he/she will be on. To achieve this, the customer may encrypt the flight number under a public key whose corresponding private key is stored only in the ADF’s trusted hardware. Thus only enclave code will be able to view the unencrypted flight number, preventing even a malicious ADF or data center operator from compromising privacy. Once the ADF is made aware of the requests (as detailed in section 2), it proceeds by decrypting the flight number and determining whether or not the relevant flight was canceled. It strips all identifying information from response, and returns either “Canceled” or “Not canceled.”

F Applications

- micro-insurance:
 - weather
 - item delivery
 - flight delay (note: something we could implement easily ourselves...)
- Ethereum and USD exchange
- stock price
- BTC exchange (with and without ADF)
- Simple derivatives
- sports betting

G Version 1

Here we define and discuss proposed extensions to the Town Crier protocol.

G.1 Request Cancellation

In order to provide recourse if the system is compromised and disabled or datagrams are delayed beyond a reasonable time, there can be a way to cancel requests for a refund. The refund must withhold a fixed fee of F_{\min} in order to ensure that malicious aborts cannot bankrupt the ADF, but the rest of the fee can be refunded at any time. If the ADF attempts to deliver a datagram for a canceled request, it will simply receive the F_{\min} needed to cover its gas costs for the attempted delivery and not deliver any data.

In order to safely handle request cancellations, we now have to store verification data on the blockchain itself. This will be considerably more expensive, but the Ethereum protocol supports it cleanly. For notational simplicity, we use four blockchain storage functions. They functions create a map from integers to arbitrary data values in the domain V .

- $\text{store} : \mathbb{N} \times V \rightarrow \emptyset$. This stores a key with an associated value in the map and returns nothing.
- $\text{load} : \mathbb{N} \rightarrow V \cup \{\perp\}$. This returns the value associated with the given key or \perp if the key is not in the map.
- $\text{storeContains} : \mathbb{N} \rightarrow \{0, 1\}$. This returns whether or not the key exists in the map.
- $\text{remove} : \mathbb{N} \rightarrow \emptyset$. This removes the key and its associated value if it is in the map and otherwise does nothing.

Table 2 describes the new \mathcal{C}_{TC} blockchain. The rest of the protocol need to change (save for calls to Deliver requiring slightly different arguments).

Using the same adversarial model, we can make the same guarantees of this new system as we did for the original Town Crier system. Even if a malicious user cancels their request just as Deliver is being called, the cancellation fee is enough to reimburse \mathcal{P}_{SGX} for any gas costs.

If we expand the adversarial model to allow for arbitrary denial of service attacks against the Town Crier system (but not the blockchain), the new Cancel functionality allows affected users to recover most of their fee with no action from the Town Crier system.

G.2 Service-level Agreements

We start by noting that a service-level agreement (SLA) is generally implemented by paying recompense if it is violated. This seems extreme if the service is not run for a profit, so thus we will assume that the costs of any SLA payments are funded by profits gained when the SLA is

\mathcal{C}_{TC} with Cancellation	
Init:	Set $\text{reqs} := \emptyset$ and $\text{reqCnt} := 0$
Request:	Upon receiving (type, callback, \$fee) from a user \mathcal{P} : If $(\$fee < F_{\min} \text{ or } \$fee > F_{\max})$ Return with no effect. Set $\text{reqID} := \text{reqCnt}$. Set $\text{reqCnt} := \text{reqCnt} + 1$. store($\text{reqID} \mapsto (\mathcal{P}, \$fee)$). Return reqID .
Deliver:	Upon receiving (reqID, data, callback) from a user \mathcal{P} : If $\mathcal{P} \neq \mathcal{P}_{SGX}$ Return with no effect. If !storeContains(reqID) Send F_{\min} to \mathcal{P}_{SGX} . Return with no further effect. $(*, \$fee) \leftarrow \text{load}(\text{reqID})$. Call callback(data) providing $\$fee - F_{\min}$ ether as the maximum gas. Send $\$fee$ ether to \mathcal{P}_{SGX} . remove(reqID).
Cancel:	Upon receiving (reqID) from a user \mathcal{P} : If !storeContains(reqID) Return with no effect. $(\mathcal{R}, \$fee) \leftarrow \text{load}(\text{reqID})$. If $\mathcal{P} \neq \mathcal{R}$ Return with no effect. Send $\$fee - F_{\min}$ to \mathcal{P} . remove(reqID).

Table 2: Definition of the \mathcal{C}_{TC} contract with cancellation.

not violated. For Town Crier, these profits can be implemented by simply increasing F_{\min} above the gas cost necessary to run Deliver. In this case, the extra money will be profit. Note that if this happens, the cancellation fee could remain simply enough to recoup gas costs and not include the profit.

In this system, an SLA could consist of a maximum amount of time before a datagram is delivered. If the user wishes to cancel a request before that time, it would be considered a voluntary cancellation and incur a cancellation fee high enough to cover gas costs of an attempted delivery. If, however, the SLA has expired before the request is canceled, then not only would a cancellation not incur a fee, the user would be returned their entire initial fee and an SLA-violation recompense. This could be a small value that would be needed to later be paid back by Town Crier system out of the profits from successfully-deliver requests in order to prevent the contract from going bankrupt.

This mechanism presents some danger if a large num-

ber of SLAs are violated at the same time and the Town Crier system is unable to provide enough funds to the contract to make all of the recompense payments. In this case, there could be a lightweight function on the contract to inform a user whether or not there is sufficient funding to make an SLA payment. This function could either before cancellation requests or it could be before a request is made. The former case would attempt to guarantee that a cancellation request right now would include an SLA payment. The latter would attempt to ensure that a new request would always have money set aside to pay for an SLA violation. Both of these utility functions may be subject to a race condition of another user making a cancellation or new request between the utility call on the actual call, thus costing an honest user money.

H Basic Protocol

H.1 A Gas-Free Basic Protocol

For simplicity, we first describe a gas-free version of our basic protocol. This basic protocol improves the strawman solution by resolving the aforementioned two issues.

Enclave-specific keys. To avoid having to verify a group signature on the blockchain, during enclave initialization, we have each enclave generate its enclave-specific key pair denoted (pk_{sgx}, sk_{sgx}) . The sk_{sgx} is retained within the enclave and used to sign the datagrams extracted from data sources during the request phase. Since Ethereum itself already verifies signatures on messages sent from users (i.e., users interact with the Ethereum blockchain through an authenticated channel), we devise a trick to *piggyback the signature verification on top of Ethereum’s already existing signature verification mechanism*. This means that the SGX enclave must sign datagrams using the [elaine: fill in name] signature scheme that is compatible with Ethereum’s signature verification. This way, we need not implement a separate signature verification in the user-defined \mathcal{C}_{TC} contract. This saves not only software engineering effort, but more importantly, gas.

To make this idea fully work, in an offline phase, a user must verify an SGX attestation vouching for its own enclave-specific public key pk_{sgx} . This pk_{sgx} is hardcoded inside the blockchain contract \mathcal{C}_{TC} . The user is responsible for making sure that this hardcoded pk_{sgx} has an appropriate SGX attestation before interacting with the \mathcal{C}_{TC} blockchain contract.

Instantiating trusted absolute clock. Since SGX’s trusted clock provides only relative time with respect to a

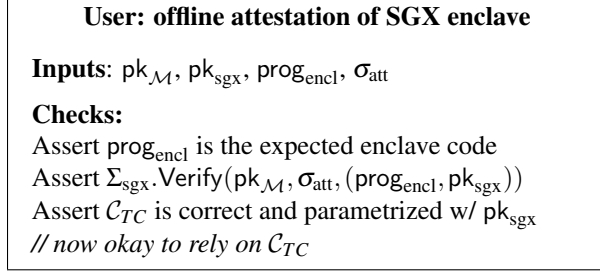


Figure 7: A user checks the Town Crier blockchain contract \mathcal{C}_{TC} , and verifies an SGX attestation of the enclave's code and its public key pk_{sgx} before entering a contract that calls \mathcal{C}_{TC} . [elaine: here we use a simplified abstraction, but the actual implementation also involves verifying the revocation list.]

reference point, we will rely on the following mechanism to realize a trusted *absolute* clock.

- **Offline calibration.** In an offline phase, a user U performs the following calibration protocol with the SGX enclave:

[elaine: this formal notation needs to be changed, it is not compatible with other formal notation.]

```

 $\mathcal{U}$ : get absolute  $T_0$  from a trusted source
 $\mathcal{U}$ : pick random nonce
 $\mathcal{U} \rightarrow \mathcal{F}_{sgx}$ : nonce
 $\mathcal{F}_{sgx} \rightarrow \mathcal{U}$ : (clock_ref,  $\Delta T_0$ , nonce)
 $\mathcal{U}$ : record clock_ref,  $\Delta T_0$ 

```

[elaine: the above assumes that an authenticated channel has been established.]

- **Online trusted absolute clock.** Whenever \mathcal{F}_{sgx} gives the relative time ΔT with respect to clock_ref, the user i) checks that clock_ref agrees with the saved reference point, and ii) computes $T_0 + \Delta T - \Delta T_0$ as the absolute time.

Formal protocol description.

H.2 Formal Guarantees

Authenticity. Roughly speaking, authenticity means that an adversary cannot convince the Town Crier blockchain contract \mathcal{C}_{TC} to accept a wrong data feed. Here a wrong data feed means any content that differs from the expected content obtained by crawling the specified url at the specified time T .

In formally defining authenticity, we assume that the user and the blockchain contract \mathcal{C}_{TC} behave honestly. Recall that the user must verify upfront the attestation σ_{att} that vouches for the enclave's public key pk_{sgx} .

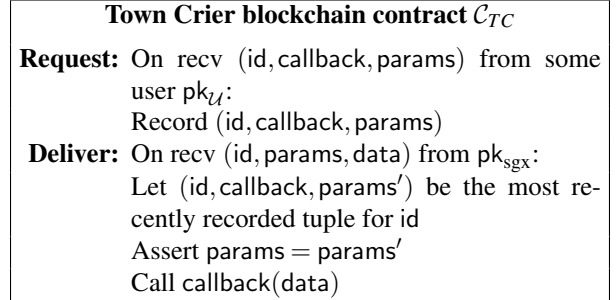


Figure 8: A simple, fee-free version of the Town Crier contract \mathcal{C}_{TC} . Note that communication with \mathcal{C}_{TC} is through an authenticated channel implemented through digital signatures (which are not explicitly expressed in our notation).

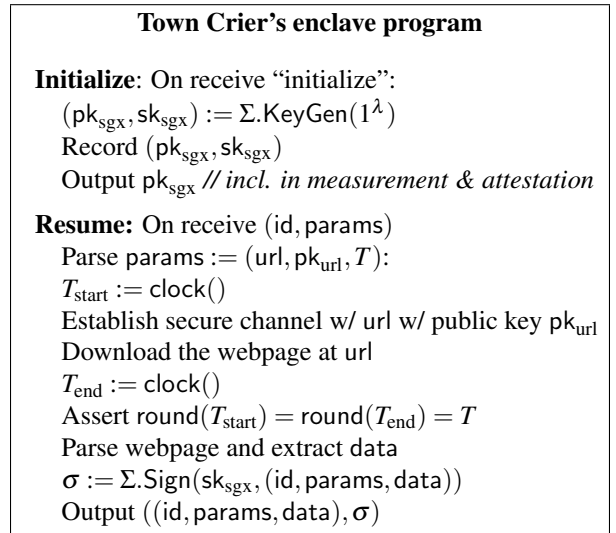


Figure 9: SGX enclave's code.

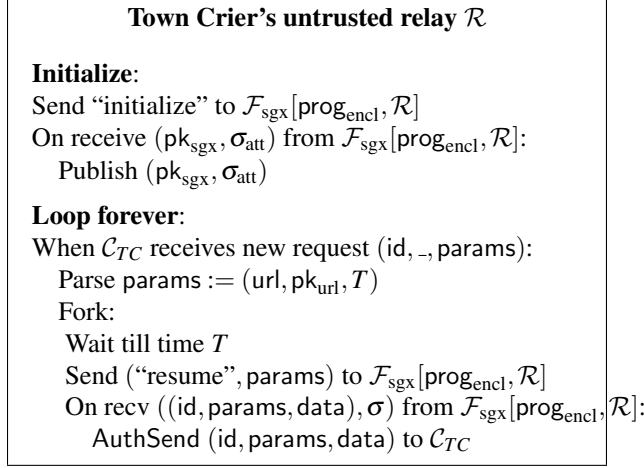


Figure 10: Town Crier untrusted relay. For simplicity, here we assume that there is only a single enclave program. When multiple data feed sources are supported, we need multiple enclaves that instantiate different parsers for different sites. In this case, the Town Crier relay also initialize all enclave instances and route the request to the correct enclave instance.

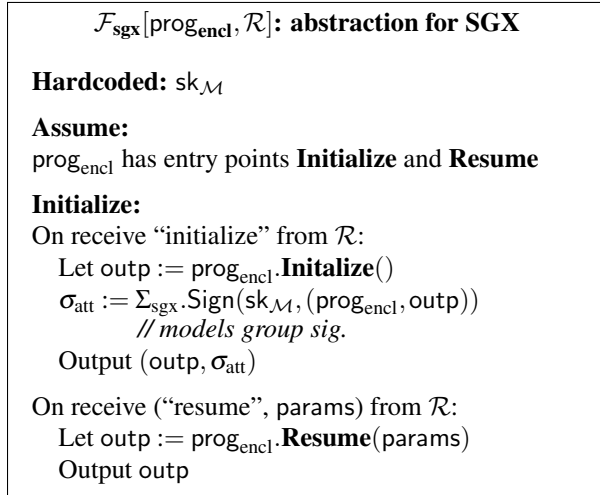


Figure 11: Formal abstraction for SGX attested execution. We adopt a similar modeling approach by Shi et al., where the SGX group signature is abstracted with a normal signature by a manufacturer key $\text{pk}_{\mathcal{M}}$. [elaine: cite our sok paper] The above functionality only models a subset of SGX features that is sufficient for our formalism.

Definition 1 (Authenticity). We say that the Town Crier protocol satisfies authenticity of data feed, if for any polynomial-time adversary that can interact arbitrarily with \mathcal{F}_{sgx} , it cannot persuade an honest verifier to accept a tuple $(\text{pk}_{\text{sgx}}, \sigma_{\text{att}}, \text{params} := (\text{url}, \text{pk}_{\text{url}}, T), \text{data}, \sigma)$ where data is not the contents of url with the public key pk_{url} at time T . More formally, for any probabilistic polynomial-time adversary \mathcal{A}

$$\Pr \left[\begin{array}{l} (\text{pk}_{\text{sgx}}, \sigma_{\text{att}}, \text{id}, \text{params}, \text{data}, \sigma) \leftarrow \mathcal{A}^{\mathcal{F}_{\text{sgx}}}(1^\lambda) : \\ (\Sigma_{\text{sgx}}.\text{Verify}(\text{pk}_{\mathcal{M}}, \sigma_{\text{att}}, (\text{prog}_{\text{encl}}, \text{pk}_{\text{sgx}})) = 1) \wedge \\ (\Sigma.\text{Verify}(\text{pk}_{\text{sgx}}, \text{id}, \text{params}, \text{data}) = 1) \wedge \\ \text{data} \neq \text{prog}_{\text{encl}}(\text{params}) \end{array} \right] \leq \text{negl}(\lambda)$$

Theorem 1 (Authenticity). Assume that Σ_{sgx} and Σ are secure signature schemes (recall that we follow Shi et al. [elaine: cite] who show how to abstractly model SGX’s group signature as a regular signature scheme under a manufacturer public key $\text{pk}_{\mathcal{M}}$), then, the above protocol achieves authenticity of data feed by Definition 1.

Proof. (sketch.) We show that if the adversary \mathcal{A} succeeds in a forgery with non-negligible probability, we can construct an adversary \mathcal{B} that can either break Σ_{sgx} or Σ with non-negligible probability. We consider two cases. The reduction \mathcal{B} will flip a random coin to guess which case it is, and if the guess is wrong, simply abort.

- Case 1: \mathcal{A} outputs a signature σ that uses the same pk_{sgx} as the SGX functionality \mathcal{F}_{sgx} . In this case, \mathcal{B} will try to break Σ . \mathcal{B} interacts with a signature challenger Ch who generates some $(\text{pk}^*, \text{sk}^*)$ pair, and gives to \mathcal{B} the public key pk^* . \mathcal{B} simulates \mathcal{F}_{sgx} by implicitly letting $\text{pk}_{\text{sgx}} := \text{pk}^*$. Whenever \mathcal{F}_{sgx} needs to sign a query, \mathcal{B} passes the signing query onto the signature challenger Ch .

Since $\text{data} \neq \text{prog}_{\text{encl}}(\text{params})$, \mathcal{B} cannot have queried Ch on a tuple of the form $(_, \text{params}, \text{data})$. Therefore, \mathcal{B} simply outputs what \mathcal{A} outputs (suppressing unnecessary terms) as the signature forgery.

- Case 2: \mathcal{A} outputs a signature σ that uses a different pk_{sgx} as the SGX functionality \mathcal{F}_{sgx} . In this case, \mathcal{B} will seek to break Σ_{sgx} . \mathcal{B} interacts with a signature challenger Ch , who generates some $(\text{pk}^*, \text{sk}^*)$ pair, and gives to \mathcal{B} the public key pk^* . \mathcal{B} simulates \mathcal{F}_{sgx} by implicitly setting $\text{pk}_{\mathcal{M}} := \text{pk}^*$. Whenever \mathcal{F}_{sgx} needs to make a signature with $\text{sk}_{\mathcal{M}}$, \mathcal{B} simply passes the signature query onto Ch . In this case, in order for \mathcal{A} to succeed, it must produce a valid signature σ_{att} for a different public key pk' . Therefore, \mathcal{B} simply outputs this as a signature forgery.

□

Town Crier blockchain contract \mathcal{C}_{TC} with fees	
Request:	On <code>recv (id, callback, params, $\Delta F_{\text{request}}$ + $\\$F_{\text{deliver}}$)</code> from some user $\text{pk}_{\mathcal{U}}$: Assert $\$F_{\text{deliver}} \in [\text{min}, \text{max}]$ Record <code>(id, callback, params, $\\$F_{\text{deliver}}$)</code> <i>// at most $\Delta F_{\text{request}}$ gas consumed</i> <i>// all remaining gas returned to $\text{pk}_{\mathcal{U}}$</i> <i>// $\\$F_{\text{request}}$ held by contract</i>
Deliver:	On <code>recv (id, params, data, $\Delta F_{\text{deliver}}$)</code> from pk_{sgx} : Let <code>(id, callback, params', $\\$F'_{\text{deliver}}$)</code> be the most recently recorded tuple for id Assert <code>params = params'</code> Assert $\$F'_{\text{deliver}} \leq \Delta F_{\text{deliver}}$ Send $\$F'_{\text{deliver}}$ to pk_{sgx} Call <code>callback(data)</code> <i>// at most $\Delta F_{\text{deliver}}$ gas consumed</i> <i>// all remaining gas returned to pk_{sgx}</i>

Figure 12: Town Crier contract \mathcal{C}_{TC} reflecting fees. $\Delta F_{\text{request}}$ denotes the gas for executing the **Request** entry point. $\Delta F_{\text{deliver}}$ denotes the gas for executing the **Deliver** entry point that includes the user-defined callback. $\$F_{\text{deliver}}$ denotes F_{deliver} amount of explicit, non-gas currency units. Essentially, the requester first pays $\$F_{\text{deliver}}$ currency units which will be used to refund the $\Delta F_{\text{deliver}}$ amount of gas that pk_{sgx} will need to put in to call the **Deliver** entry point.

I Extensions

I.1 Handling Transaction Fees

To mitigate potential Denial-of-Service (DoS) attacks, Ethereum employs a fee mechanism, referred to as “gas”, where the submitter of a transaction (that invokes an entry point in the contract) pays a transaction fee roughly proportional to the execution time of the corresponding entry point.

Our basic Town Crier implements a policy where the requester pays for all gas needed and Town Crier in effect pays nothing. The concrete instantiation is described in Figure [elaine: refer]. Since pk_{sgx} has to invoke the **Deliver** entry point, it has to advance a gas payment $\Delta F_{\text{deliver}}$. This amount will be entirely refunded through money deposited in the contract by the requester.

I.2 Support for Cancellation